

Document Classification Using Decision Trees

Sarah Morris

Decision Trees To Solve a Mystery in History

We previously used clustering techniques to see if we could determine the authorship of the disputed federalist papers. Now, we will use decision trees to see if they can provide us with further insight into the authorship of the papers.

Data Preparation

#importing dataset

```
papers = read.csv("~/Downloads/week4_resources (2)/fedPapers85_fromClass.csv")
```

#remove filename

```
papers = papers[,-2]
```

```
head(papers)
```

```
## author a all also an and any are as at be been
```

```
## 1 dispt 0.280 0.052 0.009 0.096 0.358 0.026 0.131 0.122 0.017 0.411 0.026
```

```
## 2 dispt 0.177 0.063 0.013 0.038 0.393 0.063 0.051 0.139 0.114 0.393 0.165
```

```
## 3 dispt 0.339 0.090 0.008 0.030 0.301 0.008 0.068 0.203 0.023 0.474 0.015
```

```
## 4 dispt 0.270 0.024 0.016 0.024 0.262 0.056 0.064 0.111 0.056 0.365 0.127
```

```
## 5 dispt 0.303 0.054 0.027 0.034 0.404 0.040 0.128 0.148 0.013 0.344 0.047
```

```
## 6 dispt 0.245 0.059 0.007 0.067 0.282 0.052 0.111 0.252 0.015 0.297 0.030
```

```
## but by can do down even every for. from had has have her
```

```
## 1 0.009 0.140 0.035 0.026 0.000 0.009 0.044 0.096 0.044 0.035 0.017 0.044 0
```

```
## 2 0.000 0.139 0.000 0.013 0.000 0.025 0.000 0.076 0.101 0.101 0.013 0.152 0
```

```
## 3 0.038 0.173 0.023 0.000 0.008 0.015 0.023 0.098 0.053 0.008 0.015 0.023 0
```

```
## 4 0.032 0.167 0.056 0.000 0.000 0.024 0.040 0.103 0.079 0.016 0.024 0.143 0
```

```
## 5 0.061 0.209 0.088 0.000 0.000 0.020 0.027 0.141 0.074 0.000 0.054 0.047 0
```

```
## 6 0.037 0.186 0.000 0.000 0.007 0.007 0.007 0.067 0.096 0.022 0.015 0.119 0
```

```
## his if. in. into is it its may more must my no not
```

```
## 1 0.017 0.000 0.262 0.009 0.157 0.175 0.070 0.035 0.026 0.026 0 0.035 0.114
```

```
## 2 0.000 0.025 0.291 0.025 0.038 0.127 0.038 0.038 0.000 0.013 0 0.000 0.127
```

```
## 3 0.000 0.023 0.308 0.038 0.150 0.173 0.030 0.120 0.038 0.083 0 0.030 0.068
```

```
## 4 0.024 0.040 0.238 0.008 0.151 0.222 0.048 0.056 0.056 0.071 0 0.032 0.087
```

```

## 5 0.020 0.034 0.263 0.013 0.189 0.108 0.013 0.047 0.067 0.013 0 0.047 0.128
## 6 0.067 0.030 0.401 0.037 0.260 0.156 0.015 0.074 0.045 0.015 0 0.059 0.134
## now of on one only or our shall should so some such than
## 1 0 0.900 0.140 0.026 0.035 0.096 0.017 0.017 0.017 0.035 0.009 0.026 0.009
## 2 0 0.747 0.139 0.025 0.000 0.114 0.000 0.000 0.013 0.013 0.063 0.000 0.000
## 3 0 0.858 0.150 0.030 0.023 0.060 0.000 0.008 0.068 0.038 0.030 0.045 0.023
## 4 0 0.802 0.143 0.032 0.048 0.064 0.016 0.016 0.032 0.040 0.024 0.008 0.000
## 5 0 0.869 0.054 0.047 0.027 0.081 0.027 0.000 0.000 0.027 0.067 0.027 0.047
## 6 0 0.876 0.141 0.052 0.022 0.074 0.030 0.015 0.030 0.007 0.045 0.015 0.030
## that the their then there things this to up upon was were what
## 1 0.184 1.425 0.114 0.000 0.009 0.009 0.044 0.507 0 0.000 0.009 0.017 0.000
## 2 0.152 1.254 0.165 0.000 0.000 0.000 0.051 0.355 0 0.013 0.051 0.000 0.000
## 3 0.188 1.490 0.053 0.015 0.015 0.000 0.075 0.361 0 0.000 0.008 0.015 0.008
## 4 0.238 1.326 0.071 0.008 0.000 0.000 0.103 0.532 0 0.000 0.087 0.079 0.008
## 5 0.162 1.193 0.027 0.007 0.007 0.000 0.094 0.485 0 0.000 0.027 0.020 0.020
## 6 0.208 1.469 0.089 0.007 0.007 0.000 0.126 0.445 0 0.000 0.007 0.030 0.015
## when which who will with would your
## 1 0.009 0.175 0.044 0.009 0.087 0.192 0
## 2 0.000 0.114 0.038 0.089 0.063 0.139 0
## 3 0.000 0.105 0.008 0.173 0.045 0.068 0
## 4 0.024 0.167 0.000 0.079 0.079 0.064 0
## 5 0.007 0.155 0.027 0.168 0.074 0.040 0
## 6 0.037 0.186 0.045 0.111 0.089 0.037 0

```

Since we know that Jay is not the author of these papers, we can remove any papers written by him from this dataset. We can also remove papers written jointly by them (HM)

```

#remove Jay papers
papers=papers[papers$author!="Jay",]
#remove HM papers
papers=papers[papers$author!="HM",]
#check to ensure data is clean
head(papers)

```

```

## author a all also an and any are as at be been
## 1 dispt 0.280 0.052 0.009 0.096 0.358 0.026 0.131 0.122 0.017 0.411 0.026
## 2 dispt 0.177 0.063 0.013 0.038 0.393 0.063 0.051 0.139 0.114 0.393 0.165
## 3 dispt 0.339 0.090 0.008 0.030 0.301 0.008 0.068 0.203 0.023 0.474 0.015
## 4 dispt 0.270 0.024 0.016 0.024 0.262 0.056 0.064 0.111 0.056 0.365 0.127
## 5 dispt 0.303 0.054 0.027 0.034 0.404 0.040 0.128 0.148 0.013 0.344 0.047
## 6 dispt 0.245 0.059 0.007 0.067 0.282 0.052 0.111 0.252 0.015 0.297 0.030
## but by can do down even every for. from had has have her
## 1 0.009 0.140 0.035 0.026 0.000 0.009 0.044 0.096 0.044 0.035 0.017 0.044 0
## 2 0.000 0.139 0.000 0.013 0.000 0.025 0.000 0.076 0.101 0.101 0.013 0.152 0
## 3 0.038 0.173 0.023 0.000 0.008 0.015 0.023 0.098 0.053 0.008 0.015 0.023 0
## 4 0.032 0.167 0.056 0.000 0.000 0.024 0.040 0.103 0.079 0.016 0.024 0.143 0
## 5 0.061 0.209 0.088 0.000 0.000 0.020 0.027 0.141 0.074 0.000 0.054 0.047 0
## 6 0.037 0.186 0.000 0.000 0.007 0.007 0.007 0.067 0.096 0.022 0.015 0.119 0
## his if. in. into is it its may more must my no not
## 1 0.017 0.000 0.262 0.009 0.157 0.175 0.070 0.035 0.026 0.026 0 0.035 0.114
## 2 0.000 0.025 0.291 0.025 0.038 0.127 0.038 0.038 0.000 0.013 0 0.000 0.127
## 3 0.000 0.023 0.308 0.038 0.150 0.173 0.030 0.120 0.038 0.083 0 0.030 0.068
## 4 0.024 0.040 0.238 0.008 0.151 0.222 0.048 0.056 0.056 0.071 0 0.032 0.087
## 5 0.020 0.034 0.263 0.013 0.189 0.108 0.013 0.047 0.067 0.013 0 0.047 0.128
## 6 0.067 0.030 0.401 0.037 0.260 0.156 0.015 0.074 0.045 0.015 0 0.059 0.134
## now of on one only or our shall should so some such than
## 1 0 0.900 0.140 0.026 0.035 0.096 0.017 0.017 0.017 0.035 0.009 0.026 0.009
## 2 0 0.747 0.139 0.025 0.000 0.114 0.000 0.000 0.013 0.013 0.063 0.000 0.000
## 3 0 0.858 0.150 0.030 0.023 0.060 0.000 0.008 0.068 0.038 0.030 0.045 0.023
## 4 0 0.802 0.143 0.032 0.048 0.064 0.016 0.016 0.032 0.040 0.024 0.008 0.000
## 5 0 0.869 0.054 0.047 0.027 0.081 0.027 0.000 0.000 0.027 0.067 0.027 0.047
## 6 0 0.876 0.141 0.052 0.022 0.074 0.030 0.015 0.030 0.007 0.045 0.015 0.030
## that the their then there things this to up upon was were what
## 1 0.184 1.425 0.114 0.000 0.009 0.009 0.044 0.507 0 0.000 0.009 0.017 0.000
## 2 0.152 1.254 0.165 0.000 0.000 0.000 0.051 0.355 0 0.013 0.051 0.000 0.000
## 3 0.188 1.490 0.053 0.015 0.015 0.000 0.075 0.361 0 0.000 0.008 0.015 0.008
## 4 0.238 1.326 0.071 0.008 0.000 0.000 0.103 0.532 0 0.000 0.087 0.079 0.008

```

```
## 5 0.162 1.193 0.027 0.007 0.007 0.000 0.094 0.485 0 0.000 0.027 0.020 0.020
## 6 0.208 1.469 0.089 0.007 0.007 0.000 0.126 0.445 0 0.000 0.007 0.030 0.015
##  when which  who  will  with would your
## 1 0.009 0.175 0.044 0.009 0.087 0.192 0
## 2 0.000 0.114 0.038 0.089 0.063 0.139 0
## 3 0.000 0.105 0.008 0.173 0.045 0.068 0
## 4 0.024 0.167 0.000 0.079 0.079 0.064 0
## 5 0.007 0.155 0.027 0.168 0.074 0.040 0
## 6 0.037 0.186 0.045 0.111 0.089 0.037 0
```

Separate data into training and testing groups

```
#separate the disputed files from other
known<-papers[papers$author!="dispt",]
unknown<-papers[papers$author=="dispt",]
#use 70% of "known" papers as training set and 30% as test set
sample <- sample(c(TRUE, FALSE), nrow(known), replace=TRUE, prob=c(0.7,0.3))
train<-known[sample,]
test<-known[!sample,]
```

Now that our data is split into three parts - the disputed papers(which we will use for testing once our model is complete), 70% of our known data (for preliminary training), and 30% of our known data (for preliminary testing). Now that we have these three subsets, we can proceed to build our decision tree.

Build and Tune Decision Tree

First, load the required packages.

```
library(rpart)
library(rpart.plot)
library(e1071)
library(tree)
library(tidyverse)

## — Attaching core tidyverse packages ————— tidyverse 2.0.0 —
## ✓ dplyr    1.1.3    ✓ readr    2.1.4
## ✓ forcats  1.0.0    ✓ stringr 1.5.1
```

```
## ✓ ggplot2 3.4.4 ✓ tibble 3.2.1
## ✓ lubridate 1.9.3 ✓ tidyr 1.3.0
## ✓ purrr 1.0.2
## — Conflicts ————— tidyverse_conflicts()
—

## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag() masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

Now, create the rpart model.

Rpart model1

```
#rpart
train_rpart = rpart(author ~ ., data = train,
                     minsplit=3, cp=0)
summary(train_rpart)

## Call:
## rpart(formula = author ~ ., data = train, minsplit = 3, cp = 0)
## n= 47
##
##      CP nsplit rel error  xerror  xstd
## 1 0.9090909    0 1.0000000 1.0000000 0.2638797
## 2 0.09090909    1 0.09090909 0.1818182 0.1257997
## 3 0.00000000    2 0.00000000 0.2727273 0.1523510
##
## Variable importance
## upon there on by to an at
## 29 22 14 12 12 7 4
##
## Node number 1: 47 observations, complexity param=0.9090909
## predicted class=Hamilton expected loss=0.2340426 P(node) =1
## class counts: 36 11
```

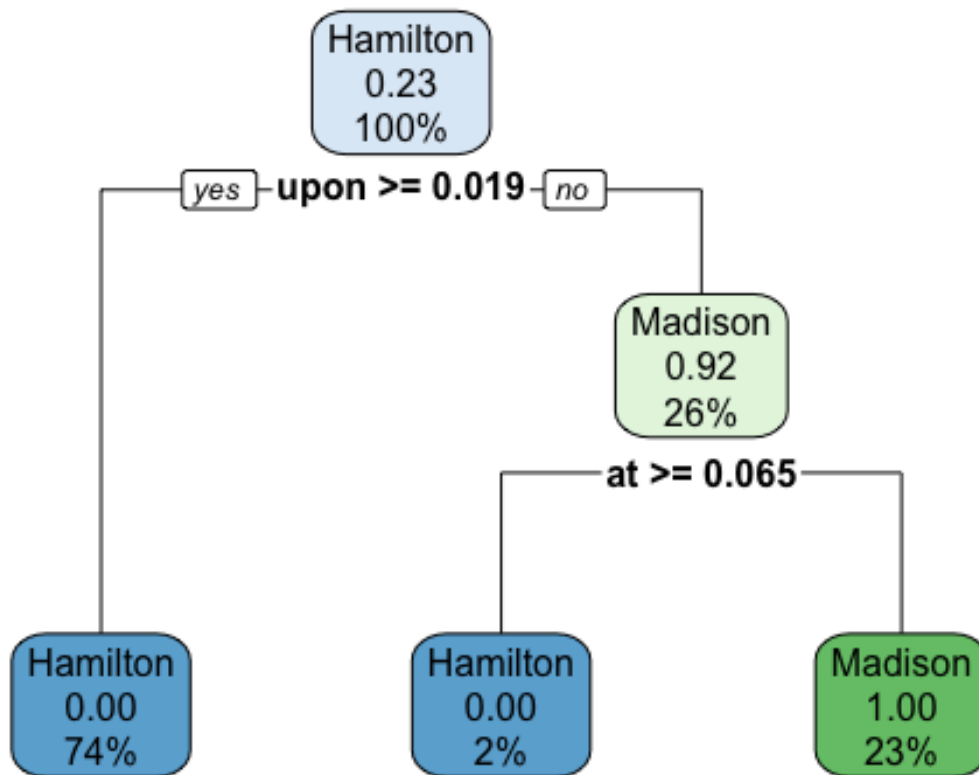
```

## probabilities: 0.766 0.234
## left son=2 (35 obs) right son=3 (12 obs)
## Primary splits:
##   upon < 0.019 to the right, improve=15.017730, (0 missing)
##   there < 0.0145 to the right, improve=10.294500, (0 missing)
##   on < 0.0825 to the left, improve= 8.579635, (0 missing)
##   to < 0.4885 to the right, improve= 7.547896, (0 missing)
##   by < 0.14 to the left, improve= 5.901064, (0 missing)
## Surrogate splits:
##   there < 0.0115 to the right, agree=0.936, adj=0.750, (0 split)
##   on < 0.0745 to the left, agree=0.872, adj=0.500, (0 split)
##   by < 0.14 to the left, agree=0.851, adj=0.417, (0 split)
##   to < 0.4885 to the right, agree=0.851, adj=0.417, (0 split)
##   an < 0.064 to the right, agree=0.809, adj=0.250, (0 split)
##
## Node number 2: 35 observations
## predicted class=Hamilton expected loss=0 P(node) =0.7446809
## class counts: 35 0
## probabilities: 1.000 0.000
##
## Node number 3: 12 observations, complexity param=0.09090909
## predicted class=Madison expected loss=0.08333333 P(node) =0.2553191
## class counts: 1 11
## probabilities: 0.083 0.917
## left son=6 (1 obs) right son=7 (11 obs)
## Primary splits:
##   at < 0.065 to the right, improve=1.833333, (0 missing)
##   be < 0.375 to the right, improve=1.833333, (0 missing)
##   even < 0.0215 to the right, improve=1.833333, (0 missing)
##   if. < 0.0025 to the left, improve=1.833333, (0 missing)
##   may < 0.026 to the left, improve=1.833333, (0 missing)
##
## Node number 6: 1 observations

```

```
## predicted class=Hamilton expected loss=0 P(node) =0.0212766
## class counts: 1 0
## probabilities: 1.000 0.000
##
## Node number 7: 11 observations
## predicted class=Madison expected loss=0 P(node) =0.2340426
## class counts: 0 11
## probabilities: 0.000 1.000
```

`rpart.plot(train_rpart)`



Rpart Model 2

Now we will build a decision tree with more nodes. We know from our tree above, that only certain variables are being used as predictors (the words upon, on, there, to, by, and should). Upon is a very strong predictor which seems to override the model. Therefore, I am removing it from this tree so we can see more nodes.

```

train_rpart2 <- rpart(author ~ on + there + to + by + should,
  data=train,
  method="class",
  control=rpart.control(minsplit=3, cp=0.002))
summary(train_rpart2)

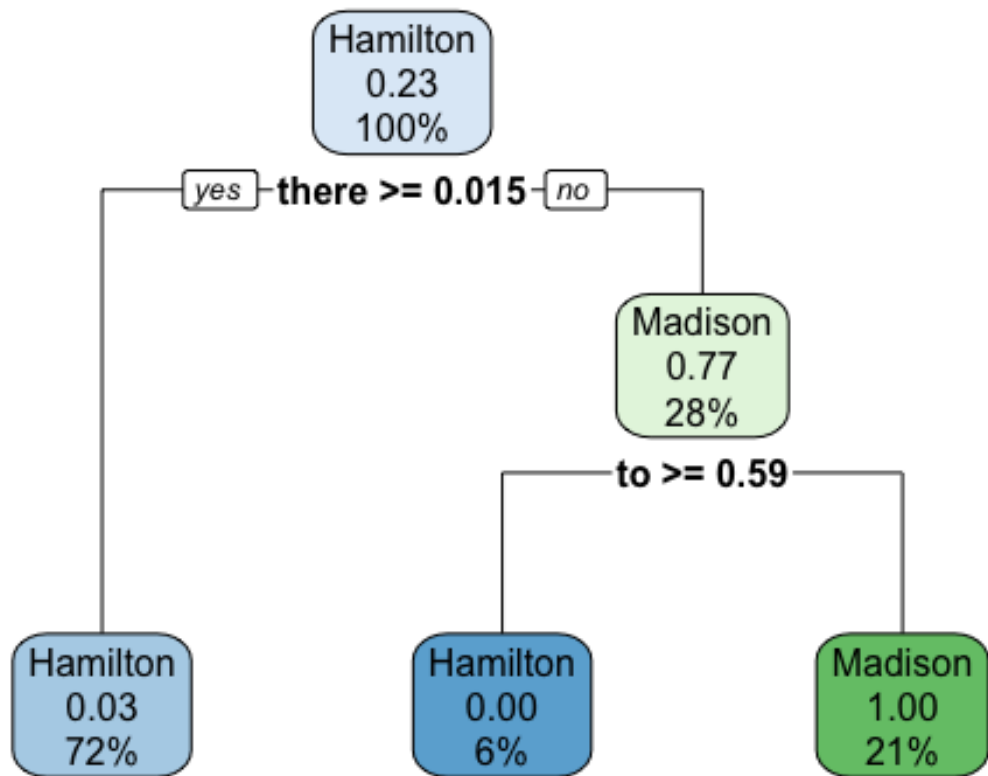
## Call:
## rpart(formula = author ~ on + there + to + by + should, data = train,
##   method = "class", control = rpart.control(minsplit = 3, cp = 0.002))
##   n= 47
##
##      CP nsplit rel error  xerror   xstd
## 1 0.6363636    0 1.00000000 1.0000000 0.2638797
## 2 0.2727273    1 0.36363636 0.6363636 0.2218898
## 3 0.0020000    2 0.09090909 0.3636364 0.1739092
##
## Variable importance
##   there    on    to    by should
##    31    23    23    12    12
##
## Node number 1: 47 observations,  complexity param=0.6363636
##   predicted class=Hamilton expected loss=0.2340426 P(node) =1
##   class counts:   36   11
##   probabilities: 0.766 0.234
##   left son=2 (34 obs) right son=3 (13 obs)
##   Primary splits:
##     there < 0.0145 to the right, improve=10.294500, (0 missing)
##     on    < 0.0825 to the left, improve= 8.579635, (0 missing)
##     to    < 0.4885 to the right, improve= 7.547896, (0 missing)
##     by    < 0.14   to the left, improve= 5.901064, (0 missing)
##     should < 0.015 to the right, improve= 2.820761, (0 missing)
##   Surrogate splits:
##     on    < 0.0915 to the left, agree=0.851, adj=0.462, (0 split)
##     to    < 0.4755 to the right, agree=0.809, adj=0.308, (0 split)

```



```
## by < 0.14 to the left, agree=0.787, adj=0.231, (0 split)
## should < 0.0445 to the left, agree=0.787, adj=0.231, (0 split)
##
## Node number 2: 34 observations
## predicted class=Hamilton expected loss=0.02941176 P(node) =0.7234043
## class counts: 33 1
## probabilities: 0.971 0.029
##
## Node number 3: 13 observations, complexity param=0.2727273
## predicted class=Madison expected loss=0.2307692 P(node) =0.2765957
## class counts: 3 10
## probabilities: 0.231 0.769
## left son=6 (3 obs) right son=7 (10 obs)
## Primary splits:
## to < 0.5875 to the right, improve=4.6153850, (0 missing)
## on < 0.0535 to the left, improve=2.7972030, (0 missing)
## should < 0.0285 to the right, improve=1.6153850, (0 missing)
## by < 0.1215 to the left, improve=1.4820510, (0 missing)
## there < 0.006 to the right, improve=0.4153846, (0 missing)
## Surrogate splits:
## on < 0.0535 to the left, agree=0.923, adj=0.667, (0 split)
## by < 0.1215 to the left, agree=0.846, adj=0.333, (0 split)
## should < 0.048 to the right, agree=0.846, adj=0.333, (0 split)
##
## Node number 6: 3 observations
## predicted class=Hamilton expected loss=0 P(node) =0.06382979
## class counts: 3 0
## probabilities: 1.000 0.000
##
## Node number 7: 10 observations
## predicted class=Madison expected loss=0 P(node) =0.212766
## class counts: 0 10
## probabilities: 0.000 1.000
```

```
#plot new train model
rpart.plot(train_rpart2)
```



Rpart model 3

Checking the variable importance before continuing to tune our decision tree:

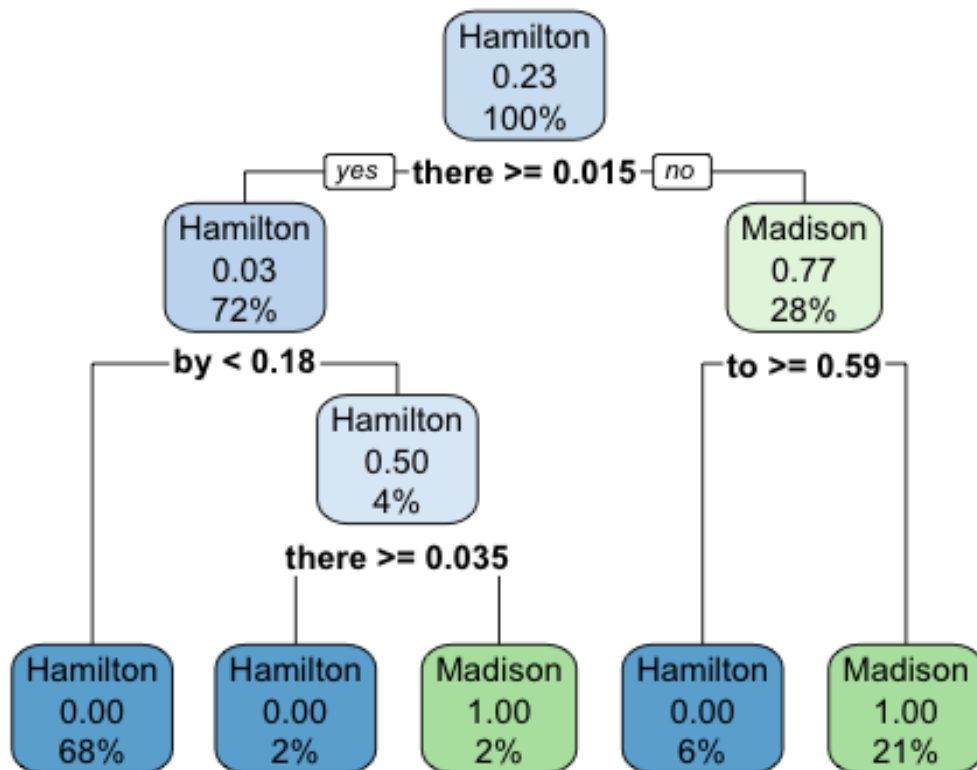
```
train_rpart$variable.importance

##   upon   there    on    by    to    an    at
## 15.017730 11.263298  7.508865  6.257388  6.257388  3.754433  1.833333
```

For the last model, we will leave out upon and on to see the new results of the decision tree.

```
train_rpart3 <- rpart(author ~ there + to + by + should,
  data=train,
  method="class",
```

```
control=rpart.control(cp=0.002,minsplit=2))
rpart.plot(train_rpart3)
```



Calculating correct classification rate:

#first model

```
ccr<-predict(train_rpart, type="class") # Model Predictions
```

```
sum(ccr != train$author)/nrow(train) # CCR
```

```
## [1] 0
```

the CCR of 0 for our first model shows that the model is 100% accurate.

#second model

```
ccr<-predict(train_rpart2, type="class") # Model Predictions
```

```
sum(ccr != train$author)/nrow(train) # CCR
```

```
## [1] 0.0212766
```

#the 0 shows that this model is also 100% correct

#third model

```
ccr<-predict(train_rpart3, type="class") # Model Predictions  
sum(ccr != train$author)/nrow(train) # CCR
```

```
## [1] 0
```

#this model is also 100% correct.

Prediction

Now we can use our models to predict on our test data:

Model 1 Prediction(train_rpart)

```
test1<-predict(train_rpart, test, type='class')  
#confusion matrix  
table(authorship=test1, true=test$author)
```

```
##      true  
## authorship Hamilton Madison  
## Hamilton      15      0  
## Madison       0      4
```

Our first model is correct the majority of the time, so we can use this on the disputed papers.

```
dispt1<-predict(train_rpart, unknown, type='class')  
table(authorship=dispt1, true=unknown$author)
```

```
##      true  
## authorship dispt  
## Hamilton    2  
## Madison     9
```

This model predicted that all of the disputed papers were written by Madison. Now, let's check our other two models to see what result they give us.

Model 2 Prediction (train_rpart2)

```
test2<-predict(train_rpart2, test, type='class')
```

```
#confusion matrix
```

```
table(authorship=test2, true=test$author)
```

```
##      true
## authorship Hamilton Madison
## Hamilton      14      0
## Madison       1      4
```

This test was slightly less accurate - 3 papers were misclassified. Let's run predict on our unknown data anyways:

```
dispt2<-predict(train_rpart2, unknown, type='class')
```

```
table(authorship=dispt2, true=unknown$author)
```

```
##      true
## authorship dispt
## Hamilton      5
## Madison       6
```

The algorithm predicted that 3 of the papers are attributed to Hamilton, though the majority of them are attributed to Madison. However, at this point, we should doubt the function of this model since it performed poorly in preliminary testing (this model left out "upon" as a deciding factor).

Model 3 Prediction

```
test3<-predict(train_rpart3, test, type='class')
```

```
#confusion matrix
```

```
table(authorship=test3, true=test$author)
```

```
##      true
## authorship Hamilton Madison
## Hamilton      14      0
## Madison       1      4
```

This model performed even worse, as it misclassified 5/24 papers (20.83% error rate)- this model left out upon and on, which were both strong predictors. We can test it on our disputed papers anyways to see the result:

```
dispt3<-predict(train_rpart3, unknown, type='class')
table(authorship=dispt3, true=unknown$author)

##      true
## authorship dispt
## Hamilton    4
## Madison     7
```

This model attributed 8 of the papers to Madison, and 3 to Hamilton.

Conclusion

Ultimately, the model 1 (which included all words in the DFM) performed the best with a 95.83% accuracy rate. This model attributed all of the papers to Madison, so I would predict with the most certainty that Madison is the author of the disputed papers. This is also supported by the results of our clustering algorithm, which also predicted that Madison was the author of the papers.

We also created 2 more models that left out the strongest predictor (upon), and then left out the two strongest predictor words (upon and on). These were understandably less accurate, but did provide us with more nodes in our decision trees relating to the other predictor words “there”, “to”, “should”, and “by”. However, we still relied on our first model for its accuracy to tell us the true authorship of the disputed papers.